

Preparing NDNP-Open-OCR for Open Source Release

September 16, 2025

This checklist outlines the essential tasks required to prepare the NDNP-Open-OCR codebase for public release, ensuring removal of internal references, elimination of sensitive information, verification of data, and compliance with licensing standards. Our goal is to complete NDNP Open OCR for open-source release by the beginning of October 2025 to ensure ample time for security review to prepare for a November open-source release.

1 Code

1.1 Remove Library of Congress-Specific References (2 pts)

- ☐ Identify and eliminate references to the loc-preservation S3 bucket or paths in the codebase (Terraform, mainly); replace with variables. Anonymize environment details, including AWS account IDs, ARNs, role names, and login URLs; provide generic configuration instructions.
- ☐ Rephrase mentions of the Library of Congress in documentation; rewrite sections to be organization-agnostic and suitable for external contributors.

1.2 Decide on Data Source for Open-Source Release (1-3 pts)

We currently use S3 as both the source and destination for OCR processing in the cloud. For the open-source release, we need to decide whether to:

- ☐ Maintain S3 as the data source and destination, providing instructions for users to set up their own S3 buckets. *This would open the door for other open-source contributors to develop other inputs*
- ☐ Generalize the data source and provide instructions for both S3 and local file system options. *This goes further than the point above, in making clear where people can build and add-on.*

1.3 Determine Redactions: Should the following be redacted or kept (1 pts)

- ☐ Gitlab CI pipeline workflows: These are used to deploy NDNP Open OCR to the LOC AWS account but may not be the case with other institutions.
- ☐ CLI workflow for kicking off batches. Should this be generalized or should it be removed since it might be too LOC-specific? Open-source users likely won't be doing the batch merging of NDNP batches like our CLI does at present. I'm thinking we should at least remove some of the sync features in the open-source release, or specify that they are LOC-specific.

1.4 Optional: Code Refactoring and Cleanup (3 pts)

If time permits, consider refactoring code for improved readability and maintainability.

2 Documentation

2.1 Draft AWS Environment Setup Recommendations (3 pts)

In the case that an organization does not have an existing IT team or AWS environment, draft a set of recommendations for setting up an appropriate AWS sandbox environment to run the NDNP-Open-OCR deployment. This should include:

- ☐ Organizational unit structure
- ☐ IAM roles and policies
- ☐ Terraform deployment (spin up and spin down)

2.2 Licensing Review (2 pts)

- ☐ Add a LICENSE file.
- ☐ Address considerations for use of American Stories models and code. Ensure it's adequately paid it's respects in public documentation.
- ☐ Verify compatibility of dependency licenses with the project's intended license; replace incompatible components if required.

2.3 Forward Looking Statements (1 pt)

I think it would be prudent to detail some ideas for further development that would inspire the open-source community to start forking the repository. This would take the form of a "Roadmap" section in the README.